

ONEPIECE OF K8S

前期准备

在开始之前，建议请先装好虚拟机，或者租一个Linux云服务器，本次课程重点任务是实践，一共给大家准备了18个实战的TASK。

DOCKER RECAP

首先我们来复习一下Docker，也让没有安装Docker的同学来安装一下环境。

1. 安装依赖：

```
sudo apt-get install \
    apt-transport-https \
    ca-certificates \
    curl \
    gnupg-agent \
    software-properties-common
```

2. 获取GPG KEY：

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-
```

3. 添加私有APT源：

```
sudo add-apt-repository \
    "deb [arch=amd64] https://download.docker.com/linux/ubuntu \
    $(lsb_release -cs) \
    stable"
```

4. 安装Docker

```
sudo apt-get install docker-ce docker-ce-cli containerd.io
```

5. 修改配置

```
vim /etc/docker/daemon.json
```

```
{
  "registry-mirrors": [
    "https://dockerhub.azk8s.cn",
    "https://reg-mirror.qiniu.com"
  ],
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "insecure-registries": ["0.0.0.0/0"],
  "storage-driver": "overlay2"
}
```

TASK 1

启动一个内有Web版2048游戏的容器：

```
docker run -d -P daocloud.io/daocloud/dao-2048
```

TASK 2

启动一个自己的镜像Registry服务：

```
docker service create --name registry --publish 5000:5000 registry:2
```

TASK 3

拉取一个Nginx镜像，并推送到自己的Registry：

```
docker pull nginx
docker tag nginx 127.0.0.1:5000/nginx
docker push 127.0.0.1:5000/nginx
```

TASK 4

这次我们来构建一个基于OpenAPI的Web APP：

1. 安装依赖：

```
pip3 install connexion flask_cors swagger-ui-bundle -i
https://mirrors.aliyun.com/pypi/simple/
```

2. 编写API配置：

```
openapi: "3.0.0"
```

```

info:
  title: Hello Pet
  version: "1.0"
servers:
  - url: http://localhost:9090/v1.0

paths:
  /pets/{pet_id}:
    get:
      description: Returns pets based on ID
      summary: Find pets by ID
      operationId: app.get_pets_by_id
      responses:
        '200':
          description: pet response
          content:
            'application/json' :
              schema:
                type: object
                $ref: '#/components/schemas/Pet'
      parameters:
        - name: pet_id
          in: path
          description: ID of pet to use
          required: true
          schema:
            type: integer
components:
  schemas:
    Pet:
      type: object
      required:
        - petType
      properties:
        petType:
          type: string
      discriminator:
        propertyName: petType
      mapping:
        dog: Dog
        cat: Cat
        lizard: Lizard
    Cat:
      allOf:
        - $ref: '#/components/schemas/Pet'
        - type: object

```

3. 用Python写后端API:

```

import connexion
from flask_cors import CORS

```

```
def get_pets_by_id(pet_id: int) -> list:
    return [{
        'name': 'lucky',
        'petType': 'dog',
        'bark': 'woof!'
    },
    {
        'name': 'kitty',
        'petType': 'cat',
        'meow': 'meow!'
    },
    {
        'name': 'jack',
        'petType': 'lizard',
        'loveRocks': True
    }][pet_id%3]

if __name__ == '__main__':
    app = connexion.FlaskApp(__name__, port=9090,
specification_dir='./')
    app.add_api('spec.yaml', arguments={'title': 'Simple Pet'})
    CORS(app.app)
    app.run()
```

4. 运行服务：

```
python app.py
```

TASK 5

基于上一个TASK，进一步构建Pet应用的镜像：

1. 写一个Dockerfile：

```
FROM python:alpine-3.8
COPY ./ /opt/app
RUN pip3 install connexion swagger-ui-bundle flask_cors -i
https://mirrors.aliyun.com/pypi/simple/
ENV PYTHONPATH /opt/app
WORKDIR /opt/app
CMD ["python", "app"]
```

2. 构建镜像：

```
docker build -t petapp -f ./Dockerfile .
```

3. 启动容器：

```
docker run petapp
```

TASK 6

继续，我们来发布Pet应用，来尝试把它发布到我们刚刚启动的私有Registry。

TASK 7

体验Portainer管理容器：

```
docker run -d -p 9000:9000 \
--name portainer --restart=always \
-v /var/run/docker.sock:/var/run/docker.sock \
portainer/portainer
```

K8S RECAP

接下来我们来复习一下Kubernetes，首先是几个基本的组件：

- etcd：etcd 用于 Kubernetes 的后端存储。所有集群数据都存储在此处，始终为您的 Kubernetes 集群的 etcd 数据提供备份计划。
- API Server：kube-apiserver对外暴露了Kubernetes API。它是的 Kubernetes 前端控制层。它被设计为水平扩展，即通过部署更多实例来缩放。
- controller manager：kube-controller-manager运行控制器，它们是处理集群中常规任务的后台线程。逻辑上，每个控制器是一个单独的进程，但为了降低复杂性，它们都被编译成独立的可执行文件，并在单个进程中运行。
- scheduler：kube-scheduler监视没有分配节点的新创建的 Pod，选择一个节点供他们运行。

TASK 8

我们现在来用kubeadm 安装 k8s。

准备工作

- 三台节点 (systemd 管理的)
- 比如debian, ubuntu, fedora, centos等
- 可以是云端节点或者虚拟机
- 分配好ip
- 配置好ssh 允许 root 登录
- 开启ipforward
- 关闭swap分区

硬件要求

- master节点 内存2核3G(最小2G)
- node节点 内存2核2G

配置 /etc/hosts

```
172.19.0.21 debian-21
172.19.0.22 debian-22
172.19.0.23 debian-23
```

安装/开启 ssh

```
apt install openssh-server
vim /etc/ssh/sshd_config
```

```
PermitRootLogin yes
```

开启 ipv4 的 forward 机制

```
vim /etc/sysctl.conf
```

```
net.ipv4.ip_forward=1
```

```
sysctl --system
```

关闭 swap

注释掉 swap 分区：

```
vim /etc/fstab
```

```
swapoff -a
```

安装 docker 运行时（略）

具体请看之前 Docker 的部分。

安装 MASTER NODE/ CONTROL PLANE

使用 kubeadm 安装 control plane：

```
apt-get update && apt-get install -y apt-transport-https
curl https://mirrors.aliyun.com/kubernetes/apt/doc/apt-key.gpg | apt-
key add -

cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
deb https://mirrors.aliyun.com/kubernetes/apt/ kubernetes-xenial main
EOF
apt-get update
apt-get install -y kubelet kubeadm kubectl
```

使用如下脚本：

```
images=(
    kube-apiserver:v1.17.0
    kube-controller-manager:v1.17.0
    kube-scheduler:v1.17.0
    kube-proxy:v1.17.0
    pause:3.1
    etcd:3.4.3-0
    coredns:1.6.5
)

for imageName in ${images[@]} ; do
    docker pull registry.cn-
hangzhou.aliyuncs.com/google_containers/$imageName
    docker tag registry.cn-
hangzhou.aliyuncs.com/google_containers/$imageName
k8s.gcr.io/$imageName
    docker rmi registry.cn-
hangzhou.aliyuncs.com/google_containers/$imageName
done
```

```
kubeadm init --apiserver-advertise-address 172.19.0.31 --pod-network-
cidr 10.30.0.0/16 --service-cidr 10.31.0.0/16
```

这个时候,似乎还是比较顺畅的,部分输出如下:

```
Your Kubernetes control-plane has initialized successfully!
```

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:

```
https://kubernetes.io/docs/concepts/cluster-administration/addons/
```

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.19.0.31:6443 --token ayly19.hyemvesfsl66wiic \  
  --discovery-token-ca-cert-hash  
sha256:b98a644c2f163ec19996599856b2d9b537ec78a0e61d920251239ec3131e5430
```

master节点运行:

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

node节点运行(在node节点搭建完成后进行):

```
kubeadm join 172.19.0.31:6443 --token ayly19.hyemvesfsl66wiic \  
  --discovery-token-ca-cert-hash  
sha256:b98a644c2f163ec19996599856b2d9b537ec78a0e61d920251239ec3131e5430
```

```
kubectl apply -f flannel.yaml
```

flannel.yaml:

```
---  
apiVersion: policy/v1beta1  
kind: PodSecurityPolicy  
metadata:  
  name: psp.flannel.unprivileged  
  annotations:  
    seccomp.security.alpha.kubernetes.io/allowedProfileNames:  
docker/default  
    seccomp.security.alpha.kubernetes.io/defaultProfileName:  
docker/default  
    apparmor.security.beta.kubernetes.io/allowedProfileNames:  
runtime/default  
    apparmor.security.beta.kubernetes.io/defaultProfileName:  
runtime/default  
spec:  
  privileged: false  
  volumes:  
  - configMap  
  - secret  
  - emptyDir  
  - hostPath  
  allowedHostPaths:  
  - pathPrefix: "/etc/cni/net.d"  
  - pathPrefix: "/etc/kube-flannel"
```



```
- pathPrefix: "/run/flannel"
readOnlyRootFilesystem: false
# Users and groups
runAsUser:
  rule: RunAsAny
supplementalGroups:
  rule: RunAsAny
fsGroup:
  rule: RunAsAny
# Privilege Escalation
allowPrivilegeEscalation: false
defaultAllowPrivilegeEscalation: false
# Capabilities
allowedCapabilities: ['NET_ADMIN']
defaultAddCapabilities: []
requiredDropCapabilities: []
# Host namespaces
hostPID: false
hostIPC: false
hostNetwork: true
hostPorts:
- min: 0
  max: 65535
# SELinux
seLinux:
  # SELinux is unused in CaaSP
  rule: 'RunAsAny'
---
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: flannel
rules:
- apiGroups: ['extensions']
  resources: ['podsecuritypolicies']
  verbs: ['use']
  resourceNames: ['psp.flannel.unprivileged']
- apiGroups:
  - ""
  resources:
  - pods
  verbs:
  - get
- apiGroups:
  - ""
  resources:
  - nodes
  verbs:
  - list
  - watch
- apiGroups:
  - ""
```

```
    resources:
      - nodes/status
    verbs:
      - patch
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1beta1
metadata:
  name: flannel
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: flannel
subjects:
- kind: ServiceAccount
  name: flannel
  namespace: kube-system
---
apiVersion: v1
kind: ServiceAccount
metadata:
  name: flannel
  namespace: kube-system
---
kind: ConfigMap
apiVersion: v1
metadata:
  name: kube-flannel-cfg
  namespace: kube-system
  labels:
    tier: node
    app: flannel
data:
  cni-conf.json: |
    {
      "name": "cbr0",
      "cniVersion": "0.3.1",
      "plugins": [
        {
          "type": "flannel",
          "delegate": {
            "hairpinMode": true,
            "isDefaultGateway": true
          }
        },
        {
          "type": "portmap",
          "capabilities": {
            "portMappings": true
          }
        }
      ]
    }
]
```

```
    }
    net-conf.json: |
    {
        "Network": "10.30.0.0/16",
        "Backend": {
            "Type": "vxlan"
        }
    }
}
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: kube-flannel-ds-amd64
  namespace: kube-system
  labels:
    tier: node
    app: flannel
spec:
  selector:
    matchLabels:
      app: flannel
  template:
    metadata:
      labels:
        tier: node
        app: flannel
    spec:
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: beta.kubernetes.io/os
                    operator: In
                    values:
                      - linux
                  - key: beta.kubernetes.io/arch
                    operator: In
                    values:
                      - amd64
      hostNetwork: true
      tolerations:
        - operator: Exists
          effect: NoSchedule
      serviceAccountName: flannel
      initContainers:
        - name: install-cni
          image: quay.io/coreos/flannel:v0.11.0-amd64
          command:
            - cp
          args:
            - -f
```

```

- /etc/kube-flannel/cni-conf.json
- /etc/cni/net.d/10-flannel.conflist
volumeMounts:
- name: cni
  mountPath: /etc/cni/net.d
- name: flannel-cfg
  mountPath: /etc/kube-flannel/
containers:
- name: kube-flannel
  image: quay.io/coreos/flannel:v0.11.0-amd64
  command:
  - /opt/bin/flanneld
  args:
  - --ip-masq
  - --kube-subnet-mgr
  resources:
    requests:
      cpu: "100m"
      memory: "50Mi"
    limits:
      cpu: "100m"
      memory: "50Mi"
  securityContext:
    privileged: false
    capabilities:
      add: ["NET_ADMIN"]
  env:
  - name: POD_NAME
    valueFrom:
      fieldRef:
        fieldPath: metadata.name
  - name: POD_NAMESPACE
    valueFrom:
      fieldRef:
        fieldPath: metadata.namespace
  volumeMounts:
  - name: run
    mountPath: /run/flannel
  - name: flannel-cfg
    mountPath: /etc/kube-flannel/
volumes:
- name: run
  hostPath:
    path: /run/flannel
- name: cni
  hostPath:
    path: /etc/cni/net.d
- name: flannel-cfg
  configMap:
    name: kube-flannel-cfg

```

apiVersion: apps/v1

```
kind: DaemonSet
metadata:
  name: kube-flannel-ds-arm64
  namespace: kube-system
  labels:
    tier: node
    app: flannel
spec:
  selector:
    matchLabels:
      app: flannel
  template:
    metadata:
      labels:
        tier: node
        app: flannel
    spec:
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: beta.kubernetes.io/os
                    operator: In
                    values:
                      - linux
                  - key: beta.kubernetes.io/arch
                    operator: In
                    values:
                      - arm64
      hostNetwork: true
      tolerations:
        - operator: Exists
          effect: NoSchedule
      serviceAccountName: flannel
      initContainers:
        - name: install-cni
          image: quay.io/coreos/flannel:v0.11.0-arm64
          command:
            - cp
          args:
            - -f
            - /etc/kube-flannel/cni-conf.json
            - /etc/cni/net.d/10-flannel.conflist
          volumeMounts:
            - name: cni
              mountPath: /etc/cni/net.d
            - name: flannel-cfg
              mountPath: /etc/kube-flannel/
      containers:
        - name: kube-flannel
          image: quay.io/coreos/flannel:v0.11.0-arm64
```

```

    command:
    - /opt/bin/flanneld
    args:
    - --ip-masq
    - --kube-subnet-mgr
    resources:
      requests:
        cpu: "100m"
        memory: "50Mi"
      limits:
        cpu: "100m"
        memory: "50Mi"
    securityContext:
      privileged: false
      capabilities:
        add: ["NET_ADMIN"]
    env:
    - name: POD_NAME
      valueFrom:
        fieldRef:
          fieldPath: metadata.name
    - name: POD_NAMESPACE
      valueFrom:
        fieldRef:
          fieldPath: metadata.namespace
    volumeMounts:
    - name: run
      mountPath: /run/flannel
    - name: flannel-cfg
      mountPath: /etc/kube-flannel/
    volumes:
    - name: run
      hostPath:
        path: /run/flannel
    - name: cni
      hostPath:
        path: /etc/cni/net.d
    - name: flannel-cfg
      configMap:
        name: kube-flannel-cfg
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: kube-flannel-ds-arm
  namespace: kube-system
  labels:
    tier: node
    app: flannel
spec:
  selector:
    matchLabels:

```

```
    app: flannel
template:
  metadata:
    labels:
      tier: node
      app: flannel
  spec:
    affinity:
      nodeAffinity:
        requiredDuringSchedulingIgnoredDuringExecution:
          nodeSelectorTerms:
            - matchExpressions:
                - key: beta.kubernetes.io/os
                  operator: In
                  values:
                    - linux
                - key: beta.kubernetes.io/arch
                  operator: In
                  values:
                    - arm
    hostNetwork: true
    tolerations:
      - operator: Exists
        effect: NoSchedule
    serviceAccountName: flannel
    initContainers:
      - name: install-cni
        image: quay.io/coreos/flannel:v0.11.0-arm
        command:
          - cp
        args:
          - -f
          - /etc/kube-flannel/cni-conf.json
          - /etc/cni/net.d/10-flannel.conflist
        volumeMounts:
          - name: cni
            mountPath: /etc/cni/net.d
          - name: flannel-cfg
            mountPath: /etc/kube-flannel/
    containers:
      - name: kube-flannel
        image: quay.io/coreos/flannel:v0.11.0-arm
        command:
          - /opt/bin/flanneld
        args:
          - --ip-masq
          - --kube-subnet-mgr
    resources:
      requests:
        cpu: "100m"
        memory: "50Mi"
      limits:
```

```

        cpu: "100m"
        memory: "50Mi"
securityContext:
  privileged: false
  capabilities:
    add: ["NET_ADMIN"]
env:
- name: POD_NAME
  valueFrom:
    fieldRef:
      fieldPath: metadata.name
- name: POD_NAMESPACE
  valueFrom:
    fieldRef:
      fieldPath: metadata.namespace
volumeMounts:
- name: run
  mountPath: /run/flannel
- name: flannel-cfg
  mountPath: /etc/kube-flannel/
volumes:
- name: run
  hostPath:
    path: /run/flannel
- name: cni
  hostPath:
    path: /etc/cni/net.d
- name: flannel-cfg
  configMap:
    name: kube-flannel-cfg
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: kube-flannel-ds-ppc64le
  namespace: kube-system
  labels:
    tier: node
    app: flannel
spec:
  selector:
    matchLabels:
      app: flannel
  template:
    metadata:
      labels:
        tier: node
        app: flannel
    spec:
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:

```



```
nodeSelectorTerms:
  - matchExpressions:
    - key: beta.kubernetes.io/os
      operator: In
      values:
        - linux
    - key: beta.kubernetes.io/arch
      operator: In
      values:
        - ppc64le
hostNetwork: true
tolerations:
- operator: Exists
  effect: NoSchedule
serviceAccountName: flannel
initContainers:
- name: install-cni
  image: quay.io/coreos/flannel:v0.11.0-ppc64le
  command:
    - cp
  args:
    - -f
    - /etc/kube-flannel/cni-conf.json
    - /etc/cni/net.d/10-flannel.conflist
  volumeMounts:
    - name: cni
      mountPath: /etc/cni/net.d
    - name: flannel-cfg
      mountPath: /etc/kube-flannel/
containers:
- name: kube-flannel
  image: quay.io/coreos/flannel:v0.11.0-ppc64le
  command:
    - /opt/bin/flanneld
  args:
    - --ip-masq
    - --kube-subnet-mgr
  resources:
    requests:
      cpu: "100m"
      memory: "50Mi"
    limits:
      cpu: "100m"
      memory: "50Mi"
  securityContext:
    privileged: false
    capabilities:
      add: ["NET_ADMIN"]
  env:
    - name: POD_NAME
      valueFrom:
        fieldRef:
```

```

        fieldPath: metadata.name
      - name: POD_NAMESPACE
        valueFrom:
          fieldRef:
            fieldPath: metadata.namespace
    volumeMounts:
      - name: run
        mountPath: /run/flannel
      - name: flannel-cfg
        mountPath: /etc/kube-flannel/
    volumes:
      - name: run
        hostPath:
          path: /run/flannel
      - name: cni
        hostPath:
          path: /etc/cni/net.d
      - name: flannel-cfg
        configMap:
          name: kube-flannel-cfg
---
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: kube-flannel-ds-s390x
  namespace: kube-system
  labels:
    tier: node
    app: flannel
spec:
  selector:
    matchLabels:
      app: flannel
  template:
    metadata:
      labels:
        tier: node
        app: flannel
    spec:
      affinity:
        nodeAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
            nodeSelectorTerms:
              - matchExpressions:
                  - key: beta.kubernetes.io/os
                    operator: In
                    values:
                      - linux
              - key: beta.kubernetes.io/arch
                    operator: In
                    values:
                      - s390x

```

```
hostNetwork: true
tolerations:
- operator: Exists
  effect: NoSchedule
serviceAccountName: flannel
initContainers:
- name: install-cni
  image: quay.io/coreos/flannel:v0.11.0-s390x
  command:
  - cp
  args:
  - -f
  - /etc/kube-flannel/cni-conf.json
  - /etc/cni/net.d/10-flannel.conflist
  volumeMounts:
  - name: cni
    mountPath: /etc/cni/net.d
  - name: flannel-cfg
    mountPath: /etc/kube-flannel/
containers:
- name: kube-flannel
  image: quay.io/coreos/flannel:v0.11.0-s390x
  command:
  - /opt/bin/flanneld
  args:
  - --ip-masq
  - --kube-subnet-mgr
  resources:
    requests:
      cpu: "100m"
      memory: "50Mi"
    limits:
      cpu: "100m"
      memory: "50Mi"
  securityContext:
    privileged: false
    capabilities:
      add: ["NET_ADMIN"]
  env:
  - name: POD_NAME
    valueFrom:
      fieldRef:
        fieldPath: metadata.name
  - name: POD_NAMESPACE
    valueFrom:
      fieldRef:
        fieldPath: metadata.namespace
  volumeMounts:
  - name: run
    mountPath: /run/flannel
  - name: flannel-cfg
    mountPath: /etc/kube-flannel/
```

```
volumes:
  - name: run
    hostPath:
      path: /run/flannel
  - name: cni
    hostPath:
      path: /etc/cni/net.d
  - name: flannel-cfg
    configMap:
      name: kube-flannel-cfg
```

flannel 国内镜像:

```
docker pull quay-mirror.qiniu.com/coreos/flannel:v0.11.0
docker tag quay-mirror.qiniu.com/coreos/flannel:v0.11.0
quay.io/coreos/flannel:v0.11.0-amd64
```

部署Node节点

部分内容和master节点雷同, 比如运行时安装与配置, kubelet安装。

1. 拉取国内镜像。
2. 部署kubelet, 配置join加入集群。

TASK 9

部署petapp到k8s集群。

(扩展: 挂载在nginx后, scale到4个pod, 创建NodePortService)

kubectl运行的三种方式

1. Generators (Run, Expose)

```
kubectl run --generator=run-pod/v1 nginx --image=nginx --image-pull-policy=IfNotPresent
```

2. 用解释的方法 (Create)

```
kubectl create deployment --image=nginx nginx
```

3. 用声明的方法 (Apply)

```
# deployment.yaml

apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: nginx-deployment
```

```
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 80
```

```
kubectl apply -f https://k8s.io/examples/application/deployment.yaml
```

K8S: 存储部分

TASK 10

emptyDir

当 Pod 指定到某个节点上时，首先创建的是一个 `emptyDir` 卷，并且只要 Pod 在该节点上运行，卷就一直存在。就像它的名称表示的那样，卷最初是空的。尽管 Pod 中的容器挂载 `emptyDir` 卷的路径可能相同也可能不同，但是这些容器都可以读写 `emptyDir` 卷中相同的文件。当 Pod 因为某些原因被从节点上删除时，`emptyDir` 卷中的数据也会永久删除。

挂载一个emptyDir的pod

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pd
spec:
  containers:
    - image: k8s.gcr.io/test-webserver
      name: test-container
      volumeMounts:
        - mountPath: /cache
          name: cache-volume
  volumes:
    - name: cache-volume
      emptyDir: {}
```

TASK 11

搭建NFS (Network File System)

service端的安装

安装包:

```
sudo apt-get install nfs-kernel-server nfs-common
```

创建共享目录:

```
sudo mkdir -p /var/nfsshare/xxx  
sudo chmod -R 777 /var/nfsshare/xxx
```

在 /etc/exports 加入:

```
/var/nfsshare/xxx    xxx.xxx.xxx.*  
(rw,sync,no_root_squash,no_all_squash)
```

或者类似:

```
/var/nfsshare/xxx    *(rw,sync,no_root_squash,no_all_squash)
```

启动服务:

```
/etc/init.d/nfs-kernel-server restart
```

TASK 12

pvc与pv的绑定

hostpath 挂载:

```
kind: PersistentVolume  
metadata:  
  name: data  
  namespace: data  
spec:  
  capacity:  
    storage: 5Gi  
  accessModes:  
    - ReadWriteOnce  
  hostPath:  
    path: /home/data  
---
```

```

kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: data
  namespace: data
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  volumeName: "data"

```

使用nfs卷进行挂载:

```

apiVersion: v1
kind: PersistentVolume
metadata:
  name: data
  namespace: data
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteMany
  nfs:
    # FIXME: use the right IP
    server: 172.19.0.11
    path: "/var/nfsshare/data"
    persistentVolumeReclaimPolicy: Recycle
---
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: data
  namespace: data
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 5Gi
  volumeName: "data"

```

K8S: 网络部分

TASK 13

nginx ingress controller的安装

配置nginx ingress controller

我们这里使用独立的ingress controller, 可以使用 nginx-ingress。nginx ingress controller可以使用多种方式部署（比如pod或daemonset），也可以使用helm完成部署。

这里给出daemonset直接配置的部署方式：

```
kubectl apply -f common/ns-and-sa.yaml
kubectl apply -f common/default-server-secret.yaml
kubectl apply -f common/nginx-config.yaml
kubectl apply -f rbac/rbac.yaml
kubectl apply -f daemon-set/nginx-ingress.yaml
```

一些自定义的配置, 需要修改common/nginx-config.yaml, 比如：

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: nginx-config
  namespace: nginx-ingress
data:
  proxy-connect-timeout: "10s"
  proxy-read-timeout: "10s"
  client-max-body-size: "100m"
```

TASK 14

根据域名的不同将流量导向指定的app

apple-app

```
kind: Pod
apiVersion: v1
metadata:
  name: apple-app
  labels:
    app: apple
spec:
  containers:
    - name: apple-app
      image: hashicorp/http-echo
      args:
        - "-text=apple"
---
```



```
kind: Service
apiVersion: v1
metadata:
  name: apple-service
spec:
  selector:
    app: apple
  ports:
    - port: 5678 # Default port for image
```

banana-app

```
kind: Pod
apiVersion: v1
metadata:
  name: banana-app
  labels:
    app: banana
spec:
  containers:
    - name: banana-app
      image: hashicorp/http-echo
      args:
        - "-text=banana"
---
kind: Service
apiVersion: v1
metadata:
  name: banana-service
spec:
  selector:
    app: banana
  ports:
    - port: 5678 # Default port for image
kubectl apply -f apple.yaml
kubectl apply -f banana.yaml
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
    - http:
        paths:
          - path: /apple
            backend:
              serviceName: apple-service
              servicePort: 5678
          - path: /banana
```

```
backend:
  serviceName: banana-service
  servicePort: 5678
```

```
kubectl apply -f apple.yaml
kubectl apply -f banana.yaml
```

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - http:
      paths:
      - path: /apple
        backend:
          serviceName: apple-service
          servicePort: 5678
      - path: /banana
        backend:
          serviceName: banana-service
          servicePort: 5678
```

```
kubectl apply -f ingress.yaml
```

K8S: 监控部分

TASK 15

metric server的安装

1. 手动安装metrics-server:

```
git clone https://github.com/kubernetes-incubator/metrics-server.git
cd metrics-server/
kubectl create -f deploy/1.8+/
```

2. 拉取国内镜像

```

images=(
    metrics-server-amd64:v0.3.6
)

for imageName in ${images[@]} ; do
    docker pull registry.cn-
hangzhou.aliyuncs.com/google_containers/$imageName
    docker tag registry.cn-
hangzhou.aliyuncs.com/google_containers/$imageName
gcr.io/google_containers/$imageName
    docker rmi registry.cn-
hangzhou.aliyuncs.com/google_containers/$imageName
done

```

deploy前要修改参数，找到

```

- name: metrics-server
  image: k8s.gcr.io/metrics-server-amd64:v0.3.6
  args:
    - --cert-dir=/tmp
    - --secure-port=4443

```

注意添加args

```

--kubelet-insecure-tls=true
--kubelet-preferred-address-types=InternalIP,Hostname,ExternalIP

```

TASK 16

kube dashboard的安装

```

kubectl apply -f
https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-
beta6/aio/deploy/recommended.yaml

```

TASK 17

使用kubectl top以及kube dashboard观察k8s平台的指标

创建访问用的USER

admin-user-sa.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: admin-user
  namespace: kubernetes-dashboard
```

admin-user-crb.yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: admin-user
  namespace: kubernetes-dashboard
root@debian-11:~/app/dashboard# cat admin-user-crb.yaml
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: admin-user
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: admin-user
  namespace: kubernetes-dashboard
```

获取secret token

```
kubectl -n kubernetes-dashboard describe secret `kubectl -n kubernetes-
dashboard get secret | grep admin-user | awk '{print $1}'`
```

进阶实战

TASK 18

创建Node节点，动手配置TLS加入kube集群。